Decision Support That Stays in Its Lane

Flags a high-risk patient for review — and never writes a treatment order.

ADVISORY

— THE PROBLEM

In medicine, the autonomous system is the liability

- A model that places an order or changes a dose has left decision support behind.
- **That line is regulatory**
  it changes the approval burden entirely.
- Clinicians won't adopt a tool they can't trust to stay advisory.

`The danger isn't a missed flag. It's a system that quietly acts.`

— THE SOLUTION

It points; the clinician decides

- **It reads**
  vitals and medication context over a FHIR-style stream.
- **It flags**
  a high-risk patient earns a clinician-review advisory, with a reason.
- **It stops**
  no orders, no writeback — ever.

— HOW IT WORKS

Vitals in, advisory out — nothing more

0 **Patient observation**
HR · SpO_2 · temp · meds

1 **Feature**
abnormality · trend · contraindication

2 **Advisory**
triage · recommend-review

3 **Result**
advisory card · disclaimer

The result layer carries an advisory card and a safety disclaimer, never an order.

— THE CEILING IS THE PRODUCT

ADVISORY, enforced

`ADVISORY`	Every result requires clinician review and carries a safety disclaimer.
`No autonomous act`	The test fails if any result even looks like an order or writeback.
`Auditable`	Each advisory carries the reason it was raised.

— THE EVIDENCE

80 deterministic ticks

0

treatment orders written — enforced

`advisory only`

1

high-risk patient flagged for review

80 / 80

ticks processed

— GOING LIVE

Connect to the record, keep the guarantee

- **Swap the input edge**
  the mock observations become a live FHIR client.
- **Keep the contract**
  same typed signals, same advisory ceiling.
- **Keep the line**
  the never-writes-an-order guarantee travels with it.

“

Good clinical AI earns trust by what it refuses to do: it points at the patient who needs a human, and it stops.

The case for Demo 04

JNEOPALLIUM · DEMO 04

rakovpublic/jneopallium

It points. You decide.

Clinician-review advisories with an audit reason — and no autonomous action.

`github.com/rakovpublic/jneopallium`

ADVISORY · deterministic